

Data Structure and Algorithm Design

ME/MSc (Computer) – Pokhara University

Prepared by:

Assoc. Prof. Madan Kadariya (NCIT)

Contact: madan.kadariya@ncit.edu.np

Chapter 3:

Cache Efficient Data Structure (10 hrs)

Outline



3.5.1. Cache-Conscious BTree

3.5.2. Cache-Conscious Merge Sort,

Cache-Conscious B-Tree

- A **cache-conscious B-tree** is specifically designed to optimize memory access patterns based on CPU cache architectures.
- The goal is to minimize cache misses by aligning the B-tree nodes with cache line sizes ensuring when a no

Key Aspects:

1. Node Structure and Memory Layout:

- In a traditional B-tree, nodes optimize tree balancing and search efficiency, while in a cache-conscious B-tree, nodes are reorganized to fit into CPU cache lines (typically 64 bytes in modern CPUs), improving memory access efficiency.

2. Alignment and Packing:

- Keys are packed into nodes, ensuring that once a node is loaded into the cache, the keys that are likely to be accessed together are also brought into cache simultaneously, reducing cache misses.

3. Data Locality:

- The child pointers are also placed in the node layout to ensure that when a key is accessed, its corresponding child pointer (if necessary) is likely to be in the same cache line, improving locality.

Cache-Conscious B-Tree



- **Example:** In a traditional B-tree with a fanout of 4, each internal node holds 3 keys and 4 pointers, but nodes may not align with cache lines, causing inefficient cache usage. In a cache-conscious B-tree, nodes are packed and aligned to fit into the CPU's cache line, allowing all keys and pointers to be fetched in a single cache load, reducing cache misses and memory accesses.
- Complexity: Both B-trees and cache-conscious B-trees have the same **asymptotic time complexity** of $O(\log_b N)$ for search, insert, and delete operations. However, cache-conscious B-trees improve **real-world performance** by reducing **cache misses** and optimizing memory access, resulting in better **amortized performance** compared to traditional B-trees.
- A **Cache-Sensitive B+ Tree (CSB+ Tree)** optimizes memory access by aligning nodes to cache lines, reducing cache misses and improving performance for large datasets. It preserves the B+ tree structure, supporting efficient search, insert, and delete operations with improved cache locality.

Cache-Oblivious merge sort

Cache-Oblivious merge sort



Sequential Merge Sort: In each step, it sorts a subarray, starting with the entire array and recursing down to smaller and smaller subarrays.

- Merge sort is a classical sorting algorithm using a **divide-and-conquer** approach. The initial unsorted list is first divided in half, each half sublist is then applied the same division method until individual elements are obtained. Pairs of adjacent elements/sublists are then **merged into sorted sublists** until the one fully merged and sorted list is obtained.
- It is a stable sorting algorithm.
- The best, average and worst case time complexity of merge sort is $O(n \log n)$

MERGE-SORT(A, p, r)

```

1  if  $p \geq r$                                 // zero or one element?
2      return
3   $q = \lfloor (p + r) / 2 \rfloor$                     // midpoint of  $A[p : r]$ 
4  MERGE-SORT( $A, p, q$ )                        // recursively sort  $A[p : q]$ 
5  MERGE-SORT( $A, q + 1, r$ )                    // recursively sort  $A[q + 1 : r]$ 
6  // Merge  $A[p : q]$  and  $A[q + 1 : r]$  into  $A[p : r]$ .
7  MERGE( $A, p, q, r$ )
  
```


Cache-Oblivious merge sort

- Cache Oblivious Models are built in a way so they can be independent of constant factors, like the size of the cache memory.
- Cache-oblivious merge sort is a variant of merge sort that is designed to work efficiently across all levels of a multi-level memory hierarchy (such as CPU caches and RAM) without requiring explicit knowledge of cache sizes or block sizes. It is a **cache-oblivious algorithm**, meaning it optimizes cache performance without tuning for specific hardware parameters.

Key Concepts

- **Divide and Conquer:** Like traditional merge sort, cache-oblivious merge sort recursively divides the array into smaller subarrays until they become trivially small.
- **Recursive Merging:** Instead of a naive merge that accesses data in a non-optimal pattern, cache-oblivious merge sort ensures that subarrays fit in cache as much as possible.
- **Implicit Cache Optimization:** It leverages the structure of recursive calls to maximize **cache efficiency**, even without explicitly knowing the cache size.

Cache-Oblivious merge sort

Cache Efficiency

- The traditional merge sort incurs many **cache misses** because merging involves non-sequential memory access.
- Cache-oblivious merge sort **minimizes cache misses** by ensuring that most memory accesses are within **a single cache line** at each level of recursion.
- The algorithm works well on hierarchical memory systems (L1, L2, RAM, disk, etc.) without requiring tuning.

Time Complexity

- **Worst-case time complexity:** $O(n \log n)$ (same as classic merge sort)
- **Expected cache complexity:** $O((n/B) \log_{M/B}(n/B))$ where:
M = size of cache, B = block size of cache, n = size of input

Cache-Oblivious merge sort

Given k sorted sequences, each containing n/k elements, the merge process follows a **divide-and-conquer** strategy:

1. Split the k sequences into two equal halves:
 1. **Left group:** $k/2$ sequences
 2. **Right group:** $k/2$ sequences
2. Recursively merge the two groups.
3. Finally, merge the two resulting sequences.
 - This recursive splitting forms a **binary merge tree** with depth $O(\log k)$

Cache-Oblivious merge sort

Time Complexity Analysis

The total number of elements is n , and at each level of recursion, every element is merged exactly once. The merge process consists of $O(\log k)$ levels.

1. Number of Merge Operations per Level:

At each level, we perform $O(k)$ merging steps, and each merge operation involves **linear** time complexity.

2. Total Work Done:

- Each level processes $O(n)$ elements.
- Since we have $O(\log k)$ levels, the total time complexity is:

$$T(n, k) = O(n \log k)$$

- Thus, the **final time complexity** for merging k sorted sequences is: $O(n \log k)$



Cache Complexity Analysis

To analyze the **cache complexity**, we use the Ideal Cache Model where:

- M is the cache size.
- B is the cache block size.
- $Q(n,k)$ is the cache complexity (number of cache misses).

1. Cache Misses in a Single Merge Operation:

- A two-way merge scans two sequences sequentially, leading to $O(\frac{n}{B})$ cache misses.
- Since we use a recursive binary merge tree, the recursion depth is $O(\log k)$

2. Total Cache Misses Over All Levels:

- At each level, merging is performed sequentially, leading to: $O((\frac{n}{B}) \log k)$
- However, considering **multi-level memory hierarchy**, the cache-oblivious model refines this to: $Q(n,k) = O((\frac{n}{B}) \log_{M/B}(n/B))$
- This represents the **optimal** number of cache misses for hierarchical memory.

Cache-Oblivious merge sort

Conclusion

- The recursive merging strategy ensures that elements fit within cache as early as possible.
- **Time Complexity:** $O(n \log k)$, which is optimal.
- **Cache Complexity:** $Q(n, k) = O\left(\left(\frac{n}{B}\right) \log_{M/B}(n/B)\right)$, significantly better than a naive merge.